

1 Task 1: Heat2D on MPI

Initial remark: An attempt at solving the full task including subtasks c) and d) was made but time was not sufficient to eradicate bugs in subsections c) and d) for multigrid. Hence, the code is submitted as-is with partially functioning multigrid kernels which are deactivated by setting the `gridCount` to 1.

1.1 Implementation

The original code was refactored and extended to include the following:

- storing all grids in `double*` pointers instead of the original cumbersome `double**` pointers.
- extending `gridLevel` to hold MPI-specific information
- extending each level's grids to include halos for MPI communication
- introducing a new `world` struct for each MPI rank to hold grid-specific information
- using `MPI_Cart_create` and associated MPI calls to dynamically create a cartesian grid

1.2 Results

The CPU code was run as-is for a baseline reference. The MPI code was run for a single rank and then for 1, 2 and 4 full nodes. All outputs are given in appendix A.

1.3 Discussion

All runs exhibit the same L2Norm as required in the task specification. The strong scaling speed-up and efficiency are observed as follows.

1.3.1 Jacobi smoothing

- 24 threads: Speed-up xxx, efficiency xxx%
- 48 threads: Speed-up xxx, efficiency xxx%
- 96 threads: Speed-up xxx, efficiency xxx%

1.3.2 Residual calculation

- 24 threads: Speed-up xxx, efficiency xxx%
- 48 threads: Speed-up xxx, efficiency xxx%
- 96 threads: Speed-up xxx, efficiency xxx%

1.3.3 L2Norm calculation

- 24 threads: Speed-up xxx, efficiency xxx%
- 48 threads: Speed-up xxx, efficiency xxx%
- 96 threads: Speed-up xxx, efficiency xxx%

1.3.4 Overall

- 24 threads: Speed-up xxx, efficiency xxx%
- 48 threads: Speed-up xxx, efficiency xxx%
- 96 threads: Speed-up xxx, efficiency xxx%

Finally, looking at the overall efficiency it becomes evident that communication costs increase and efficiency decreases as we add ranks on multiple nodes as is to be expected. Optimizing communication could help with this to some extent, but ultimately the algorithm as-is suffers from having to perform an `MPI_Allreduce` for each rank to decide when to stop iterations.

2 Task 2: Communication-Tolerant Programming

2.1 Part a)

2.1.1 Implementation

Using OpenMP, the first hybrid model can be implemented by adding the line `#pragma omp parallel for collapse(3)` before calculating the Jacobi stencil. This can be done because the loops are perfectly nested and no data race is created. Further, as discussed on Piazza, the grid parameter was changed to 768 to allow for integer divisibility by the required number of ranks.

2.1.2 Results

Strangely, the L2 Norms achieved were somewhat inconsistent even though the directive of adapting the grid size to 768 given on Piazza was adhered to. This applies also to the given pure MPI code which was run otherwise unmodified. The pure MPI code was run as-is for a baseline reference on 1, 2 and 4 full nodes. The hybrid MPI/OpenMP code was run for a partial load on 1, 2 and 4 nodes and then for 1, 2 and 4 full nodes. All outputs are given in appendix B.

2.1.3 Discussion

Looking at the pure implementation, network communication becomes evident as expected when we increase from 2 to 4 nodes. For the hybrid implementation, it was attempted to unveil communication cost by running a partial load distributed across several nodes. The increasing significant `MPI_Waitall` times when the isolated MPI ranks are distributed across several nodes reflect those relevant communication cost. Running the hybrid implementation on full nodes in an attempt to reduce the requirement for communication initially shows no improvement on 2 full nodes probably as the reduced communication between ranks is compensated by increasing message size as the ranks' grid fractions are larger in this case. However, when running the hybrid model on 4 full nodes, the `MPI_Waitall` reduce by a factor of almost two as the benefits of less communication per se coincides with reduced message size, hinting that this benefit would likely increase when adding further nodes.

A Task 1 Euler outputs

Listing 1: Task 2: Collected Euler outputs.

CPU:	1 node(s) / 1 thread(s)
MPI:	1 node(s) / 1 thread(s)
MPI:	1 node(s) / 24 thread(s)
MPI:	2 node(s) / 48 thread(s)

Compute:	0.5175 s
MPI_Irecv:	0.0006 s
MPI_Isend:	0.0038 s
Packing:	0.0000 s
Unpacking:	0.0000 s
MPI_Waitall:	0.3561 s
<hr/>	
Total Time:	0.9141 s
L2 Norm:	0.9811283499

Hybrid MPI/OpenMP: 2 node(s) / 4 * 6 = 24 thread(s): PARTIAL

```
[eu-c7-104-15:08559] SETTING BINDING TO CORE
```

[illegible]

Execution Times:

```

Compute:      6.3606 s
MPI_Irecv:    0.0009 s
MPI_Isend:    0.0040 s
Packing:      0.0000 s
Unpacking:    0.0000 s
MPI_Waitall:  12.6248 s

```

Total Time:	19.1037s
L2 Norm:	0.9887095465

Hybrid MPI/OpenMP: 2 node(s) / 4 * 12 = 48 thread(s): FULL

[eu-c7-102-11:34256] SETTING BINDING TO CORE

[illegible]

Execution Times:

Compute:	3.8945 s
MPI_Irecv:	0.0022 s
MPI_Isend:	0.0034 s
Packing:	0.0000 s
Unpacking:	0.0000 s
MPI_Waitall:	22.2460 s

Total Time:	26.3027 s
L2 Norm:	0.9887095465

Pure MPI: 4 node(s) / 96 * 1 = 96 thread(s): FULL

Execution Times:

```

Execution Times:
Compute:      0.1653s
MPI_Irecv:    0.0028s
MPI_Isend:    0.0018s
Packing:      0.0000s
Unpacking:    0.0000s
MPI_Waitall:  5.6531s

```

```

Total Time:      6.0211s
L2 Norm:        0.9811237635

```

Hybrid MPI/OpenMP: 4 node(s) / 4 * 6 = 24 thread(s): PARTIAL

[eu-c7-118-08:24464] SETTING BINDING TO CORE

```
[eu-c7-101-13:17817] MCW rank 1 bound to socket 0[core 0[hwt 0-1]], socket 0[core 1[hwt 0-1]], socket
↪ 0[core 2[hwt 0-1]], socket 0[core 3[hwt 0-1]], socket 0[core 4[hwt 0-1]], socket 0[core 5[hwt
↪ 0-1]]: [BB/BB/BB/BB/BB/BB/././././././././././././././././././././././././././././././.]
[eu-c7-101-10:36177] MCW rank 3 bound to socket 0[core 0[hwt 0-1]], socket 0[core 1[hwt 0-1]], socket
↪ 0[core 2[hwt 0-1]], socket 0[core 3[hwt 0-1]], socket 0[core 4[hwt 0-1]], socket 0[core 5[hwt
↪ 0-1]]: [BB/BB/BB/BB/BB/BB/././././././././././././././././././././././././././././././.]
[eu-c7-118-08:24464] MCW rank 0 bound to socket 0[core 0[hwt 0-1]], socket 0[core 1[hwt 0-1]], socket
↪ 0[core 2[hwt 0-1]], socket 0[core 3[hwt 0-1]], socket 0[core 4[hwt 0-1]], socket 0[core 5[hwt
↪ 0-1]]: [BB/BB/BB/BB/BB/BB/././././././././././././././././././././././././././././././.]
[eu-c7-101-06:14872] MCW rank 2 bound to socket 0[core 0[hwt 0-1]], socket 0[core 1[hwt 0-1]], socket
↪ 0[core 2[hwt 0-1]], socket 0[core 3[hwt 0-1]], socket 0[core 4[hwt 0-1]], socket 0[core 5[hwt
↪ 0-1]]: [BB/BB/BB/BB/BB/BB/././././././././././././././././././././././././././././././.]
```

Execution Times:

Hybrid MPI/OpenMP: 4 node(s) / 8 * 12 = 96 thread(s): FULL

```
[eu-c7-119-05:48662] SETTING BINDING TO CORE
```

Execution Times:

Total Time:	12.4772s
L2 Norm:	0.9887095465